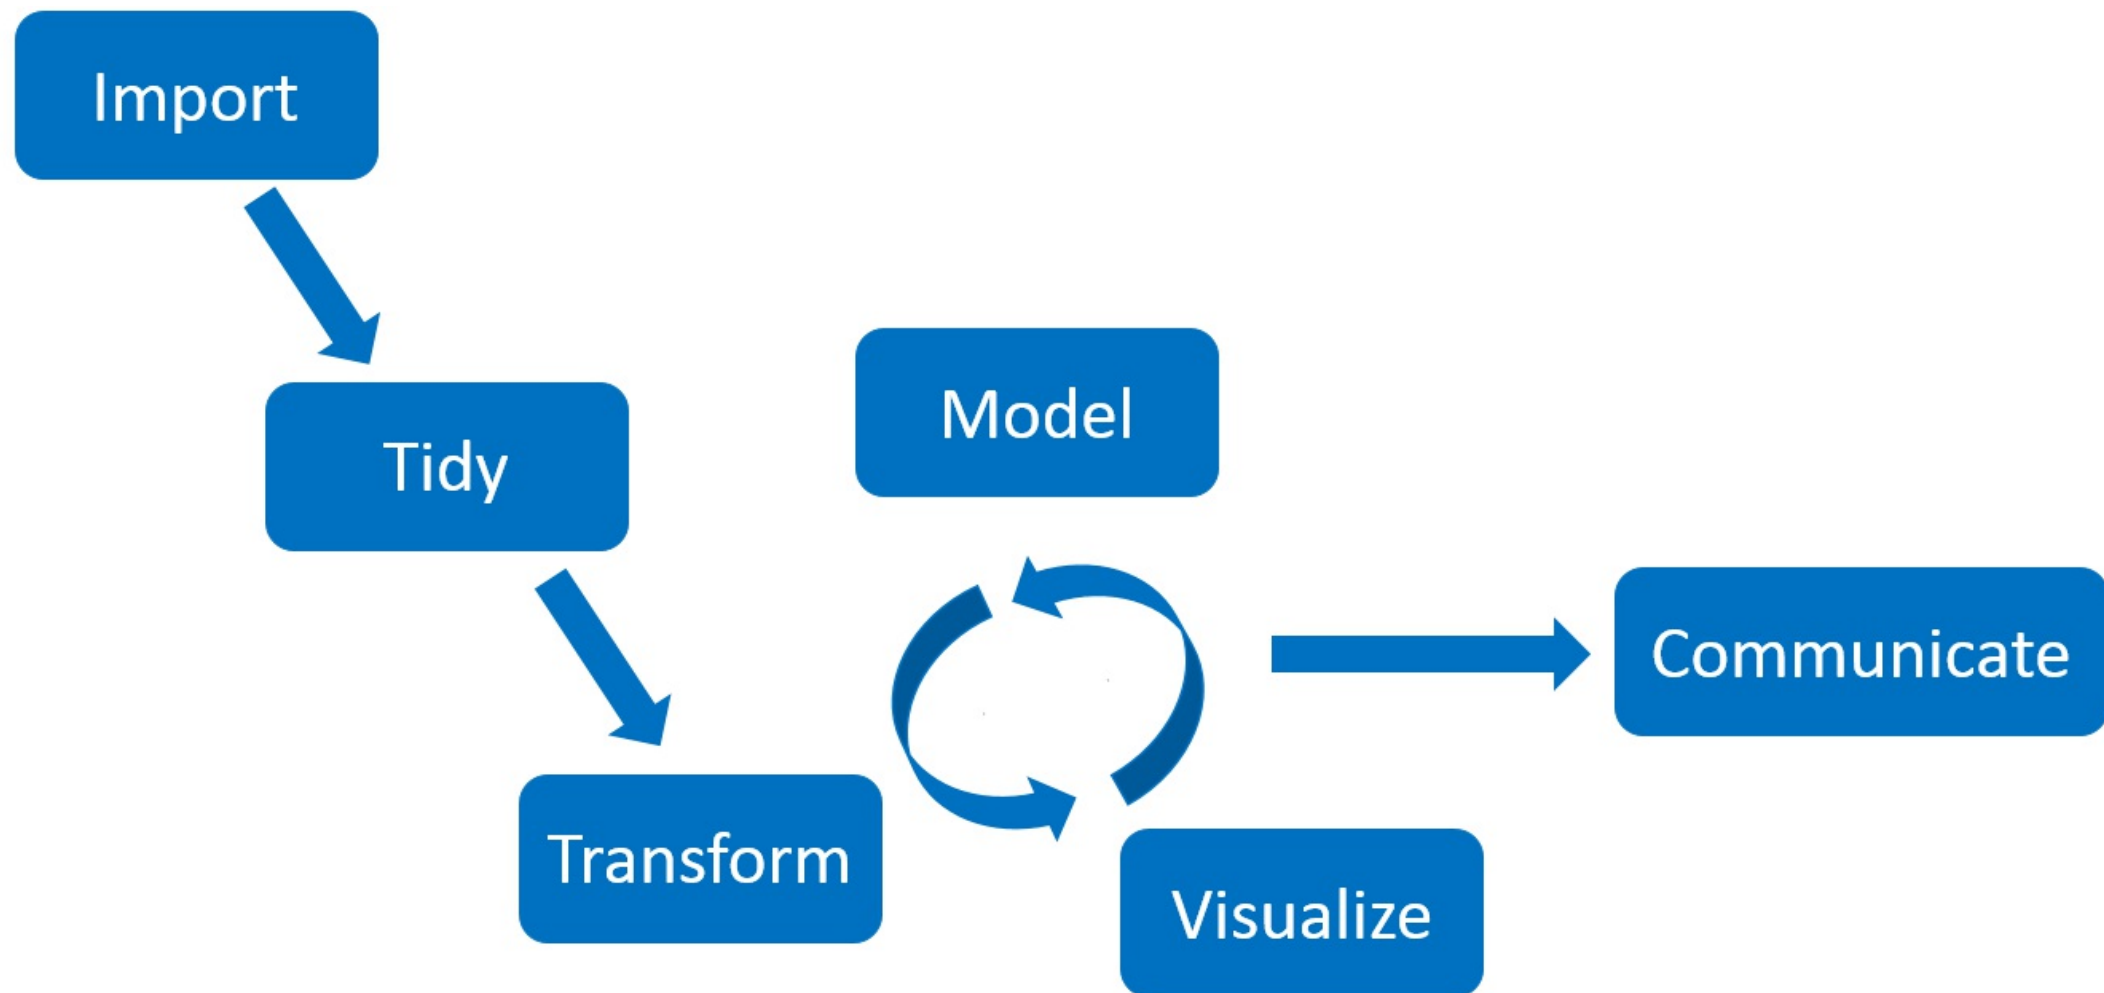# Machine learning workflow management in R
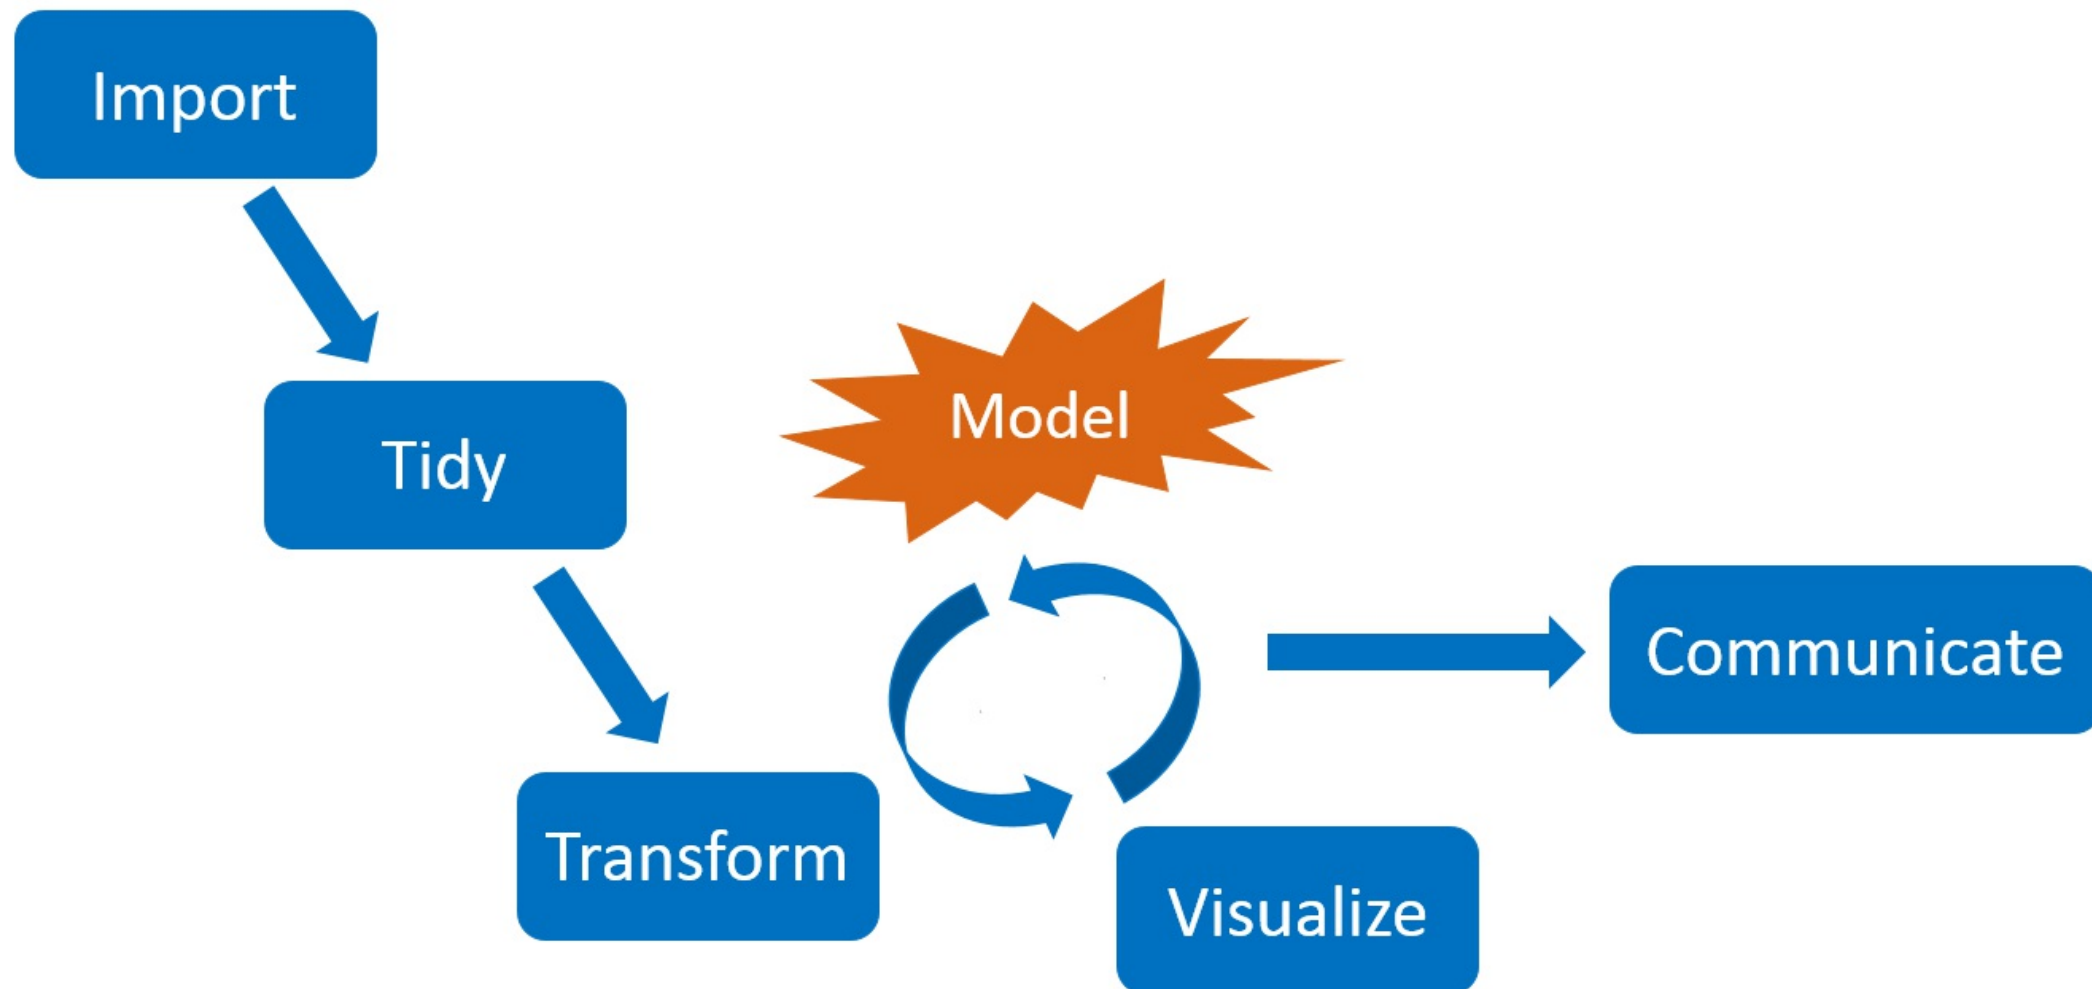

Will Landau
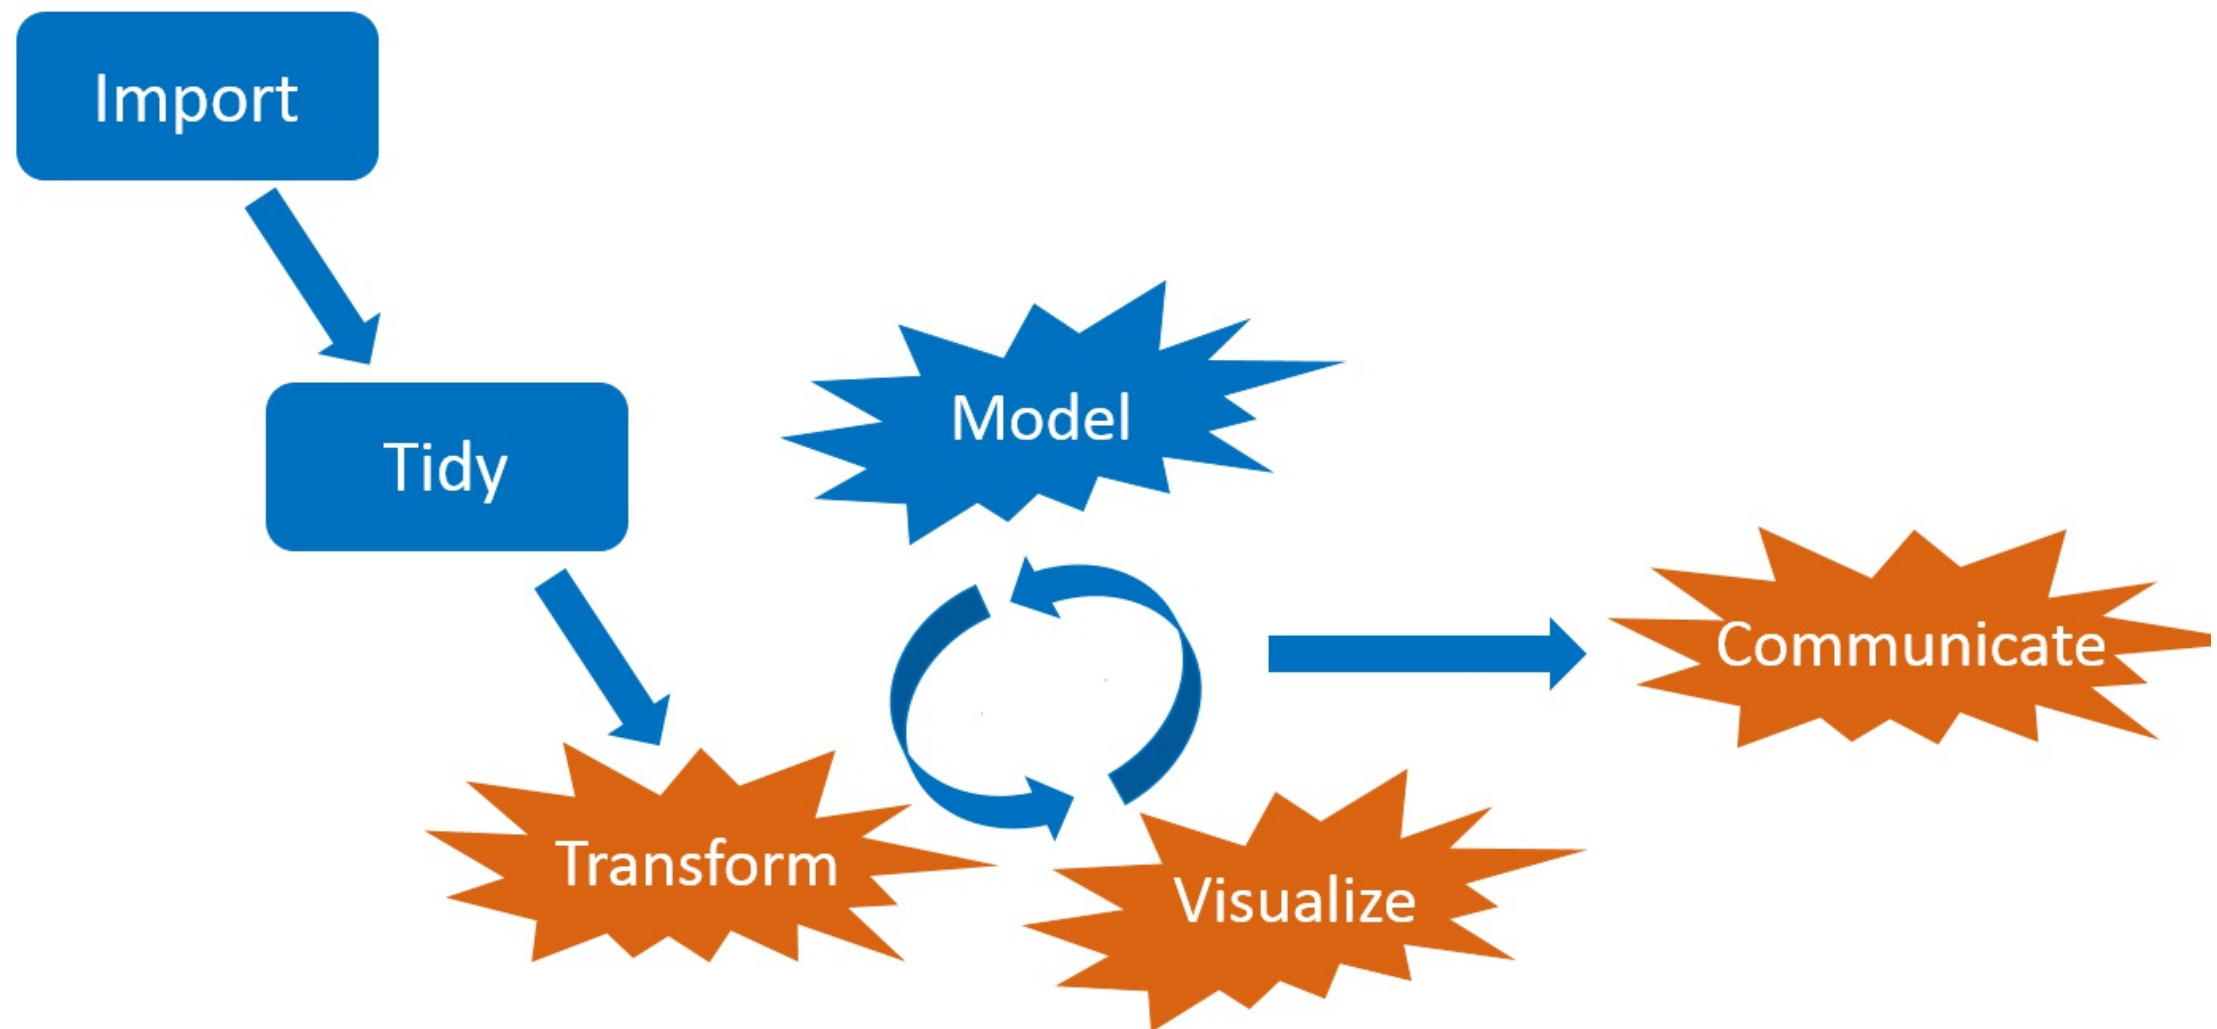
Workflows have interconnected steps.

When you change something...

...the downstream steps are no longer valid.

Do you rerun everything from scratch?

- Not if you deal with long runtimes!

https://openclipart.org/detail/275842/sisyphus-overcoming-silhouette

Do you pick and choose what to update?

- Messy.
- Prone to human error.
- Not reproducible.

https://openclipart.org/detail/216179/messy-desk

When do we face these issues?

- Long computation!
    - Clinical trial modeling and simulation
    - Subgroup identification
    - Bayesian data analysis
    - Graph-based multiple comparison procedures
    - Bayesian networks in genomics
    - PK/PD modeling (e.g. `mrgsolve`)
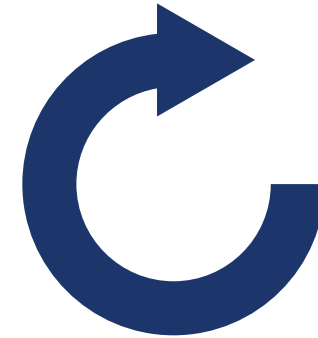   - **Deep learning**
    - …

Solution: pipeline tools

Scale up the
work you need.

Skip the
work you don't.

See evidence of
reproducibility.

- Tons exist already: github.com/pditommaso/awesome-pipeline.
- Most are language-agnostic or designed for Python or the shell.

What makes drake different?

- Aggressively designed for R.
  - Think **functions**, not script files.
  - Think **variables**, not output files.
  - Think **data frames**, not `Makefiles`.
- Major improvements in late 2018 and early 2019:
  - A domain-specific language for workflows.
  - Massive improvements in speed and memory usage.
  - Special functions and checks to safeguard reproducibility.

Example: a deep learning workflow

- Goal: predict customers who cancel their subscriptions with a telecom company.
- Data: IBM Watson Telco Customer Churn dataset.
- Workflow principles generalize to other industries.

https://openclipart.org/detail/90739/newplus, https://github.com/rstudio/keras

File structure

```
make.R
R/
├── packages.R
├── functions.R
└── plan.R
data/
└── customer_churn.csv
```

packages.R

```r
library(drake)
library(keras)
library(recipes)
library(rsample)
library(tidyverse)
library(yardstick)
```

functions.R

```r
prepare_recipe <- function(data) {
  # ...
}

define_model <- function(rec, units1, units2, act1, act2, act3) {
  # ...
}

train_model <- function(data, rec, units1, units2, act1, act2, act3) {
  # ...
}

confusion_matrix <- function(data, rec, serialized_model) {
  # ...
}

compare_models <- function(...) {
  # ...
}
```
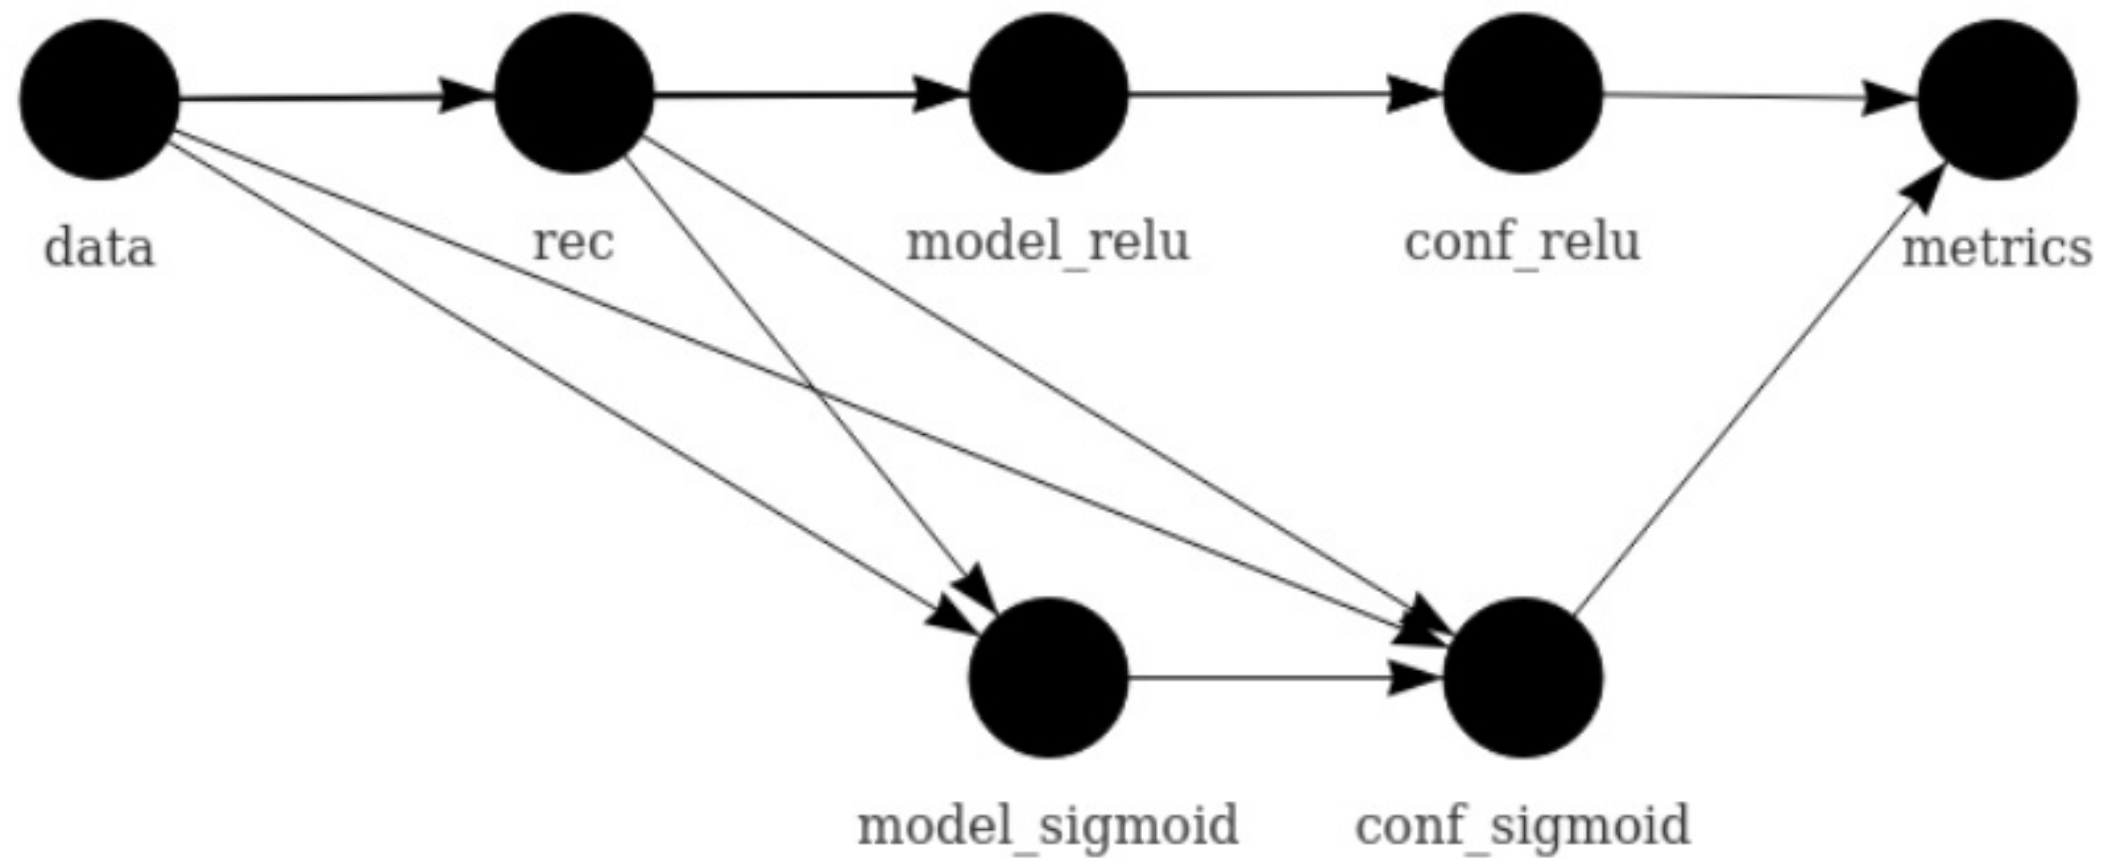
plan.R

```r
activations <- c("relu", "sigmoid")

plan <- drake_plan(
  data = read_csv(file_in("data/customer_churn.csv"), col_types = cols()) %>%
    initial_split(prop = 0.3),
  rec = prepare_recipe(data),
  model = target(
    train_model(data, rec, act1 = act),
    transform = map(act = !!activations)
  ),
  conf = target(
    confusion_matrix(data, rec, model),
    transform = map(model, .id = act)
  ),
  metrics = target(
    compare_models(conf),
    transform = combine(conf)
  )
)
```

Data frame of workflow steps

```
plan
## # A tibble: 7 x 2
##   target      command
##   <chr>       <expr>
## 1 data        read_csv(file_in("data/customer_churn.csv"), col_types = c
## 2 rec         prepare_recipe(data)
## 3 model_relu  train_model(data, rec, act1 = "relu")
## 4 model_sigmo… train_model(data, rec, act1 = "sigmoid")
## 5 conf_relu   confusion_matrix(data, rec, model_relu)
## 6 conf_sigmoid confusion_matrix(data, rec, model_sigmoid)
## 7 metrics     compare_models(conf_relu, conf_sigmoid)
```

The workflow

Run the project in make.R.

```
source("R/packages.R")
source("R/functions.R")
source("R/plan.R")

make(plan)
## target data
## target rec
## target model_relu
## target model_sigmoid
## target conf_relu
## target conf_sigmoid
## target metrics
```

Compare models.

```
readd(metrics) # See also loadd()
```

Try another activation function.

```
activations <- c("relu", "sigmoid", "softmax")

plan <- drake_plan(
  data = read_csv(file_in("data/customer_churn.csv"), col_types = cols()) %>%
    initial_split(prop = 0.3),
  rec = prepare_recipe(data),
  model = target(
    train_model(data, rec, act1 = act),
    transform = map(act = !!activations)
  ),
  conf = target(
    confusion_matrix(data, rec, model),
    transform = map(model, .id = act)
  ),
  metrics = target(
    compare_models(conf),
    transform = combine(conf)
  )
)
```

vis_drake_graph()
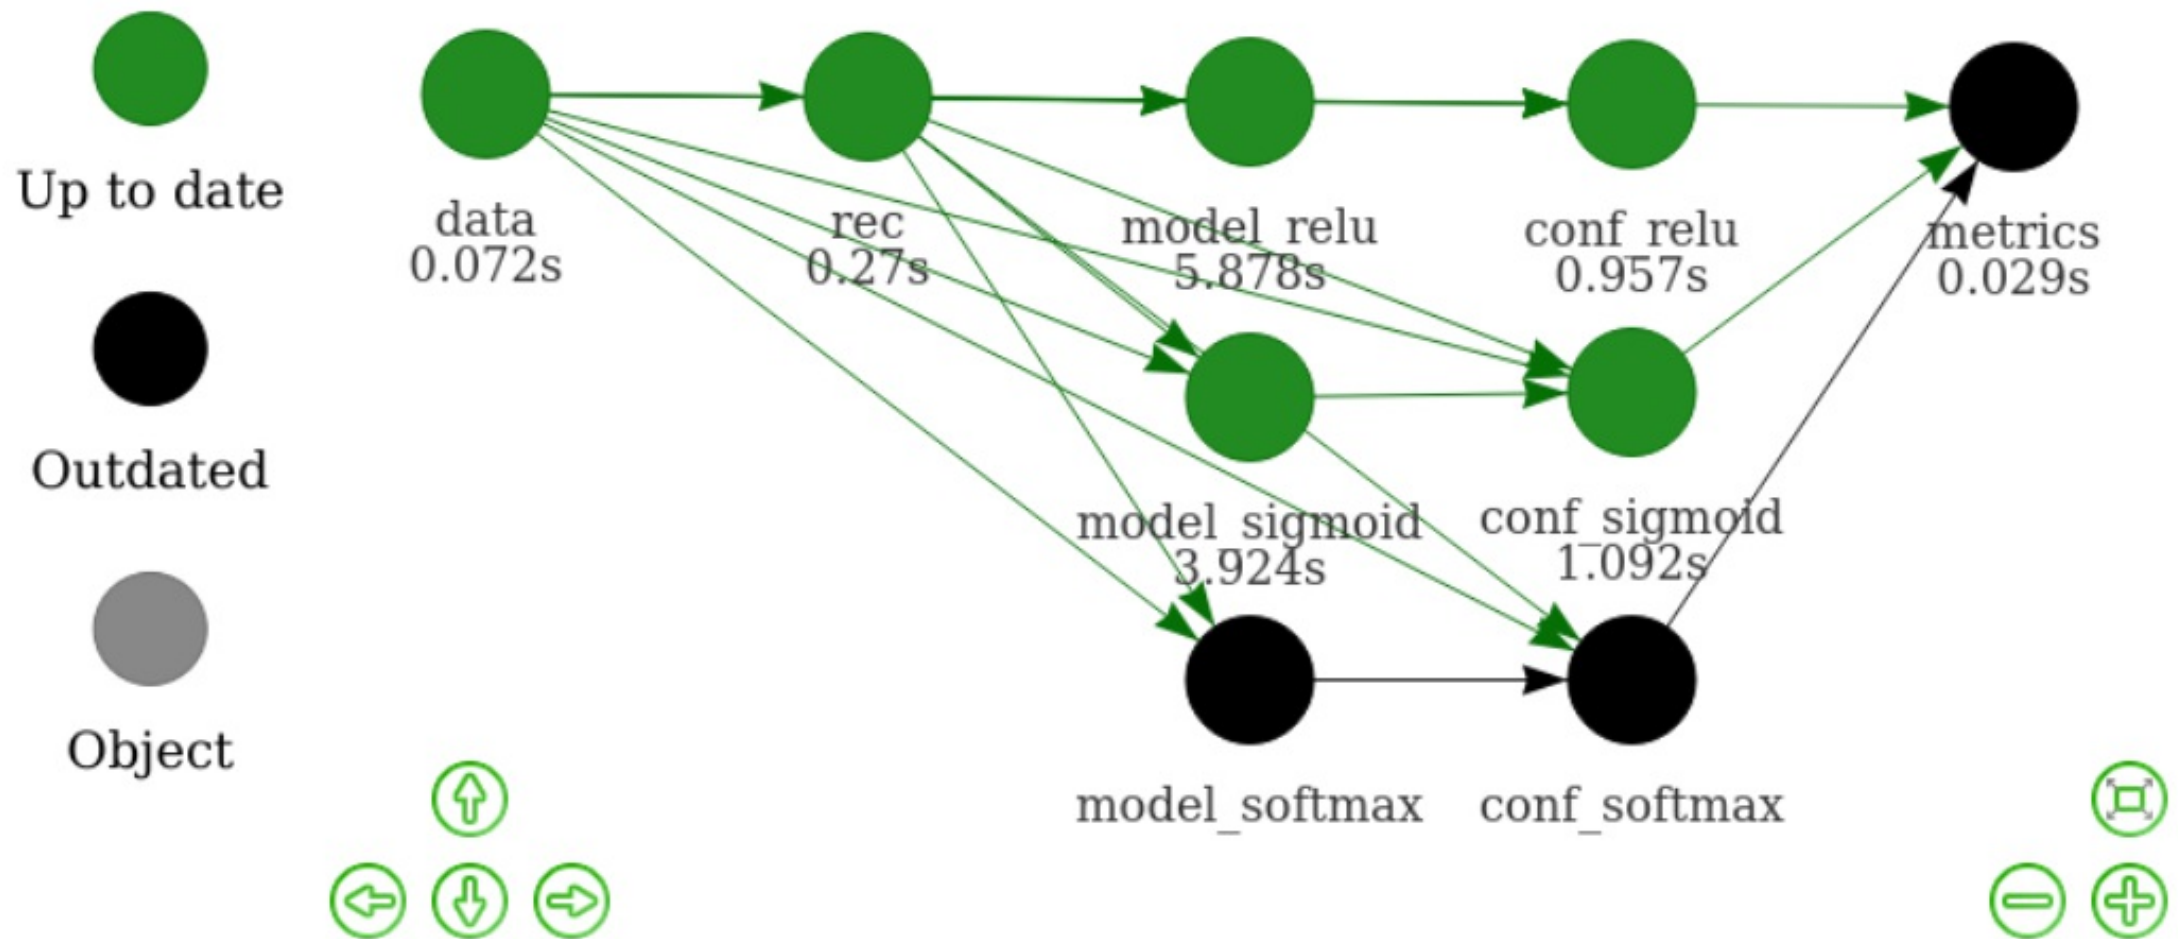
Refresh the results in make.R.

```r
source("R/packages.R")
source("R/functions.R")
source("R/plan.R") # modified

make(plan)
## target model_softmax
## target conf_softmax
## target metrics
```

Compare models.

```
readd(metrics)
```

Evidence of reproducibility

```r
source("R/packages.R")
source("R/functions.R")
source("R/plan.R")

make(plan)
## All targets are already up to date.
```

- See also `outdated()`.

History and provenance

```
history <- drake_history()
history
## # A tibble: 10 x 9
##     target  time       hash    exists command          runtime latest  prop act
##     <chr>   <chr>      <chr>   <lgl>  <chr>               <dbl> <lgl>   <dbl> <ch
##  1 conf_r… 2019-06-…  e82b0… TRUE    confusion_ma…       1.58  TRUE     NA   <NA
##  2 conf_s… 2019-06-…  e735c… TRUE    confusion_ma…       1.86  TRUE     NA   <NA
##  3 conf_s… 2019-06-…  562f7… TRUE    confusion_ma…       2.49  TRUE     NA   <NA
##  4 data    2019-06-…  ca84b… TRUE    "read_csv(fi…       0.0860 TRUE    0.3  <NA
##  5 metrics 2019-06-…  4ddc4… TRUE    compare_mode…       0.0230 FALSE    NA   <NA
##  6 metrics 2019-06-…  1a16a… TRUE    compare_mode…       0.02   TRUE     NA   <NA
##  7 model_… 2019-06-…  1b810… TRUE    "train_model…      13.5   TRUE     NA   rel
##  8 model_… 2019-06-…  4163b… TRUE    "train_model…       6.02  TRUE     NA   sig
##  9 model_… 2019-06-…  23402… TRUE    "train_model…       6.97  TRUE     NA   sof
## 10 rec     2019-06-…  7abeb… TRUE    prepare_reci…       0.164 TRUE     NA   <NA
```

High-performance computing

```
# template file with configuration
drake_hpc_template_file("slurm_clustermq.tmpl")

# Use SLURM resource manager with the template.
options(
  clustermq.scheduler = "slurm",
  clustermq.template = "slurm_clustermq.tmpl"
)

# make() is the basically the same.
make(plan, jobs = 2, parallelism = "clustermq")
```

High-performance computing

Resources

- Get drake:

```
install.packages("drake")
```

- Workshop materials:

```
remotes::install_github("wlandau/learndrake")
```

- Example code from these slides:

```
drake::drake_example("customer-churn-simple")
```

Links

- Development repository: https://github.com/ropensci/drake
- Full user manual https://ropenscilabs.github.io/drake-manual
- Reference website: https://ropensci.github.io/drake
- Code examples: https://github.com/wlandau/drake-examples
- Discuss at rOpenSci.org: https://discuss.ropensci.org

rOpenSci use cases

- Use `drake`? Share your use case at https://ropensci.org/usecases.

Thanks

- Edgar Ruiz
- example code

- Matt Dancho
- blog post

Thanks

- Maëlle Salmon
- Ben Marwick
- Julia Lowndes
- Peter Slaughter
- Jenny Bryan
- Rich FitzJohn
- Stefanie Butland

- Jarad Niemi
- Kirill Müller
- Henrik Bengtsson
- Michael Schubert
- Kendon Bell
- Miles McBain
- Patrick Schratz
- Alex Axthelm
- Jasper Clarkberg
- Tiernan Martin
- Ben Listyg
- TJ Mahr
- Ben Bond-Lamberty
- Tim Mastny
- Bill Denney
- Amanda Dobbyn
- Daniel Falster
- Rainer Krug
- Brianna McHorse
- Chan-Yub Park